

Model and notation

Point mass m on a massless leg of length ℓ , pivoting about the stance foot on a slope of inclination γ . The angle θ is measured clockwise from upward vertical, so with an ankle torque τ and moment of inertia $I = m\ell^2$ about the foot,

$$\ddot{\theta} = \frac{g}{\ell} \sin \theta + \frac{\tau}{m\ell^2}$$

At touchdown both legs are symmetric about the slope normal, which is itself tilted γ from vertical, so the event guards and impact map are

$$\theta_{TD} = \gamma + \alpha, \quad \theta^+ = \gamma - \alpha, \quad \dot{\theta}^+ = \cos(2\alpha) \dot{\theta}^-$$

Controls: $\tau \in [-0.1 mg\ell, 0.05 mg\ell]$ every timestep, and $\alpha \in [\pi/8, \pi/7]$ once per step. Parameters: $\gamma = 0.06$ rad, $\ell = 1$ m, $m = 1$ kg, $g = 9.81$ m/s².

1. Sketches

2. Region of attraction of the ankle controller

Controller

Choosing

$$\tau = m\ell^2 \ddot{\theta}_{des} - mg\ell \sin \theta, \quad \ddot{\theta}_{des} = -k_p \theta - k_d \dot{\theta}$$

cancels the gravity term exactly (feedback linearization) and leaves the closed-loop system

$$\ddot{\theta} + k_d \dot{\theta} + k_p \theta = 0$$

which is asymptotically stable at $(\theta, \dot{\theta}) = (0, 0)$ for any $k_p, k_d > 0$. This is the “invert gravity and add damping” strategy: one $-\frac{g}{\ell} \sin \theta$ of torque cancels gravity and a second supplies the restoring term. Comparing with $\ddot{\theta} + 2\zeta\omega_n \dot{\theta} + \omega_n^2 \theta = 0$ gives $k_p = \omega_n^2$ and $k_d = 2\zeta\omega_n$; I used $\zeta = 1$, i.e. $k_d = 2\sqrt{k_p}$, since critical damping is the fastest response without overshoot, and overshoot is particularly costly here because the backward torque bound ($0.05 mg\ell$) is half the forward one ($0.1 mg\ell$). Gains used: $k_p = 28.48$, $k_d = 10.67$.

The commanded torque is clipped to the permissible range. Once it saturates, gravity is no longer cancelled and the guarantee above fails, so the region of attraction must be found numerically.




Figure 1: Snapshots of the walker at mid-stance, at impact with $\alpha = \pi/8$ and with $\alpha = \pi/7$, and in the roll-back failure mode; with the corresponding state-space plot locating each snapshot and showing the event guards for both extremes of α .

Grid bounds

The RoA is found by simulating the closed loop from every point of a 90×90 grid and recording which states settle at upright ($|\theta| < 10^{-3}$ and $|\dot{\theta}| < 10^{-3}$ within 5 s).

- θ : the walker is only ever between its own stance walls, so the grid spans $\gamma \pm \pi/7 = [-0.389, 0.509]$ rad, using the largest permissible α on both sides. A run is aborted as a failure when it leaves this interval: backwards means rolling back onto the previous foot, forwards means the next foot has already landed.
- $\dot{\theta}$: with the ankle off, the states that coast exactly to upright lie on the separatrix $\dot{\theta} = -2\sqrt{g/\ell} \sin(\theta/2)$, whose magnitude is at most 1.2 rad/s over that θ range. The bounded torque can only correct small deviations from it, so $\dot{\theta} \in [-1.5, 1.5]$ rad/s covers the RoA with margin.

A static check supports these bounds: holding the pendulum at rest needs $mg\ell|\sin\theta|$, so no state with $\theta > \arcsin(0.1) \approx 0.1$ rad or $\theta < -\arcsin(0.05) \approx -0.05$ rad can be held at zero velocity by any permissible torque.

Result




Figure 2: Region of attraction of the ankle controller (90×90 grid, torque clipped to $[-0.1, 0.05] mg\ell$).

The RoA is a narrow diagonal band along the separatrix rather than a blob around the origin. With at most $0.1 mg\ell$ of torque against a gravitational torque of up to $mg\ell$, the controller can only rescue states that were already close to coasting to a stop at upright. Without any torque the set of states reaching upright is just the one-dimensional stable manifold (the incoming separatrix

branches); the bounded torque widens that curve into a band whose width is set by the torque bounds, not by the gains. Two slices:

- at $\dot{\theta} = 0$: $\theta \in [-0.036, 0.105]$ rad, asymmetric because the forward torque bound is twice the backward one, and matching the static estimate above;
- at $\theta = 0$: $\dot{\theta} \in [-0.152, 0.287]$ rad/s.

Event guard

The RoA is stored as a boolean array over the grid. The guard converts a state to grid indices and reports “inside” only if the entire 3×3 neighbourhood of grid points is inside. This conservatism matters: with a plain nearest-point lookup, states just outside the true boundary were admitted, the controller switched on, and the walker fell. The cost is that the usable band is trimmed by about one cell on each edge, which is why the RoA grid resolution matters as well (Section 4). While walking, the ankle torque is held at zero; the guard is evaluated every timestep and the controller is latched on permanently the first time it fires.




Figure 3: Event guard with α fixed at $\pi/8$, starting from $\dot{\theta}_0 = 2.0$ rad/s: the ankle stays off through three impacts, then the state enters the RoA and the controller brings the walker to rest at upright.

3. Choice of Poincaré section

With α fixed, touchdown is a convenient section because $\theta_{TD}(\gamma, \alpha)$ is constant. Once α is a control input the touchdown guard moves with the input, so it no longer defines a fixed surface and a state on it would need both $\dot{\theta}$ and the α used. I therefore use

$$\Sigma = \{(\theta, \dot{\theta}) : \theta = 0, \dot{\theta} > 0\}$$

the instant the walker passes upright moving downhill. It satisfies both requirements:

- **Transverse to the relevant flow.** The flow is $(\dot{\theta}, \ddot{\theta})$; on the line $\theta = 0$ the component normal to the line is $\dot{\theta}$, strictly positive on Σ , so orbits cross it and never run along it. Every orbit that completes a step crosses $\theta = 0$ exactly once. Orbits that stall before upright and roll back never reach it, but those are failures, not relevant flow.
- **θ constant.** Every point of Σ has $\theta = 0$, so the section state is the single number $\dot{\theta}_k$ and the step-to-step dynamics reduce to the scalar map $\dot{\theta}_{k+1} = P(\dot{\theta}_k, \alpha_k)$ with the single input α_k .

Two further advantages: Σ lies strictly between the stance walls for every permissible α ($-0.389 < 0 < 0.453$), so switching α on the section never moves a guard the walker is close to — a stance begins at $\gamma - \alpha_k$, crosses Σ , and ends at $\gamma + \alpha_{k+1}$ — and the standing equilibrium itself lies on Σ , so “has the walker reached the RoA” becomes an interval of speeds on the section.

4. Verifying the grid resolution

Table and policy

The state axis of the lookup table is $\dot{\theta}_k$ from 0 to $\sqrt{2g/\ell} = 4.43$ rad/s (Froude number $\text{Fr} = \ell\dot{\theta}^2/g = 2$); the control axis is $\alpha \in [\pi/8, \pi/7]$. Each cell is simulated from $(0, \dot{\theta}_k)$ with the ankle off until the state enters the RoA, rolls back, or returns to Σ with a new speed $\dot{\theta}_{k+1}$. Backward induction then labels rows already in the RoA as 0 steps, rows where some α reaches the RoA as 1 step, and rows whose landing speed is a $(k-1)$ -step row as k steps. Landing speeds essentially never coincide with a grid row, so they are rounded to the nearest one — which is exactly the approximation this section has to justify.

Criterion

A grid is adequate only if the real continuous walker behaves as the table predicts. I ran the full timestep simulation under the table’s policy from 150 randomly drawn starting speeds in $[0, 4.43]$ rad/s — random so that they essentially never coincide with grid rows, and with a fixed seed so every resolution faces the same test set. A grid passes if:

1. every start reaches standing (no failures);
2. no start takes more steps than the table promised for its nearest row (no “broken promises”);
3. against a 300-row reference table, the worst-case step count is unchanged and at most 5% of starts take one extra step.

Criteria 1 and 2 test correctness; criterion 3 tests optimality, since a coarse table can be correct while simply wasting steps.

Results

15 α columns throughout; the 300-row reference has a worst case of 3 steps.

rows	failed starts	broken promises	worst case	starts needing an extra step	build time [s]	passes
30	0	0	4	29.3 %	0.4	no
60	0	0	4	13.3 %	0.7	no
80	0	0	4	8.0 %	1.0	no
90	0	0	3	3.3 %	1.1	yes
100	0	0	3	4.7 %	1.3	yes
150	0	0	3	2.0 %	1.8	yes

Decision: 90 rows, the coarsest grid that passes. The slightly lower resolution of 80 rows fails on two counts: its worst case rises to 4 steps — one more than the reference achieves — and 8% of starting speeds take an extra step, against 3.3% at 90 rows. The transition in the worst-case column between 80 and 90 rows is sharp, and is the primary evidence for the choice.

Some observations:

- No resolution produced a failure or a broken promise. The neighbour margin used in the backward induction (a candidate α must also work from the rows either side, and a landing is judged by the worst of the landing row and its neighbours) makes coarse tables conservative rather than wrong: they waste steps instead of falling. Without that margin, the chosen α tends to sit at the edge of the range that works, and starts lying between rows did take more steps than promised.
- The extra-step percentage falls with resolution but not monotonically (3.3% at 90 rows, 4.7% at 100). Extra steps occur only for starting speeds near a boundary between step counts, and how many of the 150 test speeds land in those narrow bands depends on where the grid rows happen to fall. The bands shrink with refinement but never vanish, so a strict “zero extra steps” criterion would not converge at any practical resolution — hence the 5% tolerance.

RoA grid resolution. The RoA grid resolution also matters, for a different reason. With a 40×40 RoA grid, a band of section speeds around 0.2 rad/s was wrongly labelled as unable to stand: the conservative 3×3 guard trims about a cell from each edge of an already thin band, and these orbits pass close to that edge. Refining the RoA grid removed the artifact, and the RoA reported here uses 90×90 . An energy check confirms it was an artifact — a walker crossing Σ at 0.2 rad/s leaves its first impact essentially on the separatrix, which lies inside the RoA.

5. Trajectories from a three-step initial condition




Figure 4: Fewest-steps policy from $\dot{\theta}_0 = 3.63$ rad/s: three steps, commanding $\alpha = 0.421, 0.433$ and 0.437 rad. Triangles mark heelstrikes, dotted lines the impact maps, the green marker the standing equilibrium.

From $\dot{\theta}_0 = 3.63$ rad/s the fewest-steps policy stands after **3 steps**. Each impact multiplies the speed by $\cos(2\alpha)$; after the third impact the orbit lies inside the RoA band, the guard fires, and the ankle controller brings the walker to rest at upright.

The opposite question — how long the walker can keep walking from this state and still reach the RoA — comes from the same table, replacing the minimisation in the backward induction with a longest-path search in which any row that can no longer reach the RoA is excluded. From $\dot{\theta}_0 = 3.63$ rad/s the walker can take at most **5 steps** before entering the RoA, and it does so by commanding $\alpha = \pi/8$ at every crossing. That is the expected answer: $\alpha = \pi/8$ minimises the impact loss (a fraction $\cos^2(2\alpha)$ of the kinetic energy is retained, 50% at $\pi/8$ against 39% at $\pi/7$), so the walker sheds energy as slowly as possible. The fact that the longest path is finite also shows that no permissible α sustains a limit cycle on this slope: even the shortest stride loses more energy at impact than the slope returns, so the walker cannot walk indefinitely.

6. Steps to standstill as a function of initial condition

steps to standstill	initial speeds $\dot{\theta}_0$ [rad/s]
0	0.00 – 0.15
1	0.20 – 1.64

steps to standstill	initial speeds $\dot{\theta}_0$ [rad/s]
2	1.69 – 2.74
3	2.79 – 4.43

Every initial condition up to Froude number 2 is brought to standstill in at most **3 steps**. Holding α fixed at $\pi/8$ instead requires up to 5 steps over the same range (Section 5), so the benefit of choosing α appears at the fast end: from $\dot{\theta}_0 = 2.0$ rad/s a fixed $\pi/8$ also takes 3 steps, but from 3.63 rad/s the policy still takes only 3 where a fixed $\pi/8$ takes 5.

Within each region the commanded α rises with speed and saturates at $\pi/7$ at the top of the region: a longer stride retains only $\cos^2(2\alpha)$ of the kinetic energy at impact, so it brakes harder, and the fastest speeds in a region need the hardest brake to finish in that many steps. Where several strides work, the middle of the working range is commanded, which produces the flat stretches near $\alpha \approx 0.42$. At low speed the policy uses short strides, because a long stride there leaves too little energy to vault over the new foot and the walker rolls back.

Side Notes

- `models/inverted_pendulum_walker.py` — dynamics, event guard, impact map, parameters.
- `assignment_2.py` — ankle controller, RoA grid search and guard, step map, lookup table, backward induction, longest-path search, resolution study, and all figures in this report. Set `RUN_RESOLUTION_STUDY = True` to reproduce the table in Section 4.
- `phase_portrait.py` — stance-phase portrait with the Poincaré section.




Figure 5: Longest walk from the same initial condition: five steps, commanding the shortest permissible stride $\alpha = \pi/8 = 0.393$ rad at every crossing.




Figure 6: Initial section speeds coloured by the fewest steps needed to reach the RoA; dots show the first α commanded by the policy.